

Getting process statistics and modifying system call XV6

Implemented in XV6-RISCV in DOCKER (in Mac OS - M1)

Code logic:

File: *test.c*

This file is created in the user directory of xv6. Here, a new process is created using *fork()*. Child process: *exec()* replaces the memory address space of the ‘test’ process with the program passed as its parameter. A new array *argv2[]* is used to store the command line arguments from index 1 (I.e. only arguments followed by test) from *argv* and passed to *exec()*. If any error occurs and *exec()* fails, it returns to this program and appropriate message is printed and program exits with an error.

```
/xv6-riscv/user/test.c
```

```
14
15     pid = fork();
16 v     if (pid == 0) {
17         // Child process
18         exec(argv[1], argv2);
19         printf("Failed to execute %s\n", argv[1]);
20         exit(1);
21 v     } else if (pid > 0) {
22         printf("%s pid %d \n", argv[1], pid);
23         // Parent process
24         waittext(&wt, &rt, &st);
25 v     } else {
26         printf("Failed to fork %s\n", argv[1]);
27     }
28
29     exit(0);
30 }
31 |
```

The parent process calls *waittext()* to wait for the child process to complete successfully. *waittext()* is extended code of standard *wait()* call, which is mentioned further in this document.

File: *defs.h*

```
// proc.c
int      cpuid(void);
void    exit(int);
int      fork(void);
int      growproc(int);
void    proc_mapstacks(pagetable_t);
pagetable_t  proc_pagetable(struct proc *);
void    proc_freetable(pagetable_t, uint64);
int      kill(int);
int      killed(struct proc*);
void    setkilled(struct proc*);
struct cpu* mycpu(void);
struct cpu* getmycpu(void);
struct proc* myproc();
void    procinit(void);
void    scheduler(void) __attribute__((noreturn));
void    sched(void);
void    sleep(void*, struct spinlock*);
void    userinit(void);
int      wait(uint64);
void    wakeup(void*);
void    yield(void);
int      either_copyout(int user_dst, uint64 dst, void *src, uint64 len);
int      either_copyin(void *dst, int user_src, uint64 src, uint64 len);
void    procdump(void);
int      waittext(uint64, uint64, uint64);
int      ps(void);
```

File: *proc.h*

Declared variables for process creation time (*ctime*), end time (*etime*), run time (*rtime*) and total time (*total_time*) in *proc* structure.

```
// Per-process state
struct proc {
    struct spinlock lock;

    // p->lock must be held when using these:
    enum proctype state;           // Process state
    void *chan;                   // If non-zero, sleeping on chan
    int killed;                  // If non-zero, have been killed
    int xstate;                  // Exit status to be returned to parent's wait
    int pid;                     // Process ID

    // wait_lock must be held when using this:
    struct proc *parent;          // Parent process

    // these are private to the process, so p->lock need not be held.
    uint64 kstack;                // Virtual address of kernel stack
    uint64 sz;                    // Size of process memory (bytes)
    pagetable_t pagetable;        // User page table
    struct trapframe *trapframe;  // data page for trampoline.S
    struct context context;       // swtch() here to run process
    struct file *file[NFILE];     // Open files
    struct inode *cwd;            // Current directory
    char name[16];               // Process name (debugging)

    int ctime, etime, rtime;
    uint total_time;
};
```

File *syscall.c*

File: *syscall.h*

File: *usys.pl*

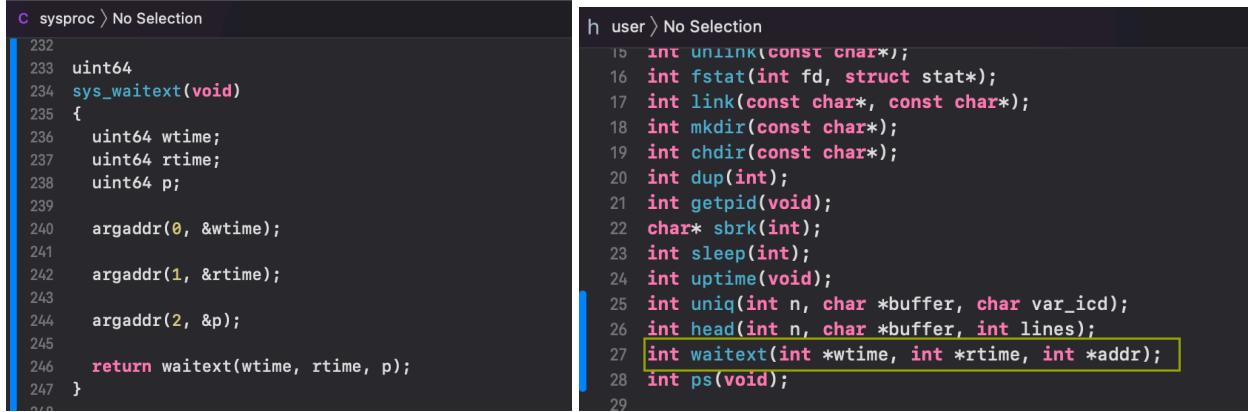
```
C syscall > No Selection
102 extern uint64 sys_mkdir(void);
103 extern uint64 sys_close(void);
104 extern uint64 sys_uniq(void);
105 extern uint64 sys_head(void);
106 extern uint64 sys_waitext(void);
107 extern uint64 sys_ps(void);
108
109 // An array mapping syscall numbers from syscall.
110 // to the function that handles the system call.
111 static uint64 (*syscalls[]) (void) = {
112     [SYS_fork] sys_fork,
113     [SYS_exit] sys_exit,
114     [SYS_wait] sys_wait,
115     [SYS_pipe] sys_pipe,
116     [SYS_read] sys_read,
117     [SYS_kill] sys_kill,
118     [SYS_exec] sys_exec,
119     [SYS_fstat] sys_fstat,
120     [SYS_chdir] sys_chdir,
121     [SYS_dup] sys_dup,
122     [SYS_getpid] sys_getpid,
123     [SYS_sbrk] sys_sbrk,
124     [SYS_sleep] sys_sleep,
125     [SYS_uptime] sys_uptime,
126     [SYS_open] sys_open,
127     [SYS_write] sys_write,
128     [SYS_mknod] sys_mknod,
129     [SYS_unlink] sys_unlink,
130     [SYS_link] sys_link,
131     [SYS_mkdir] sys_mkdir,
132     [SYS_close] sys_close,
133     [SYS_uniq] sys_uniq,
134     [SYS_head] sys_head,
135     [SYS_waitext] sys_waitext,
136     [SYS_ps] sys_ps,
137 };
138
```

```
h syscall > No Selection
1 // System call numbers
2 #define SYS_fork 1
3 #define SYS_exit 2
4 #define SYS_wait 3
5 #define SYS_pipe 4
6 #define SYS_read 5
7 #define SYS_kill 6
8 #define SYS_exec 7
9 #define SYS_fstat 8
10 #define SYS_chdir 9
11 #define SYS_dup 10
12 #define SYS_getpid 11
13 #define SYS_sbrk 12
14 #define SYS_sleep 13
15 #define SYS_uptime 14
16 #define SYS_open 15
17 #define SYS_write 16
18 #define SYS_mknod 17
19 #define SYS_unlink 18
20 #define SYS_link 19
21 #define SYS_mkdir 20
22 #define SYS_close 21
23 #define SYS_uniq 22
24 #define SYS_head 23
25 #define SYS_waitext 24
26 #define SYS_ps 25
27
```

```
P- usys > No Selection
25 entry("KILL");
26 entry("exec");
27 entry("open");
28 entry("mknod");
29 entry("unlink");
30 entry("fstat");
31 entry("link");
32 entry("mkdir");
33 entry("chdir");
34 entry("dup");
35 entry("getpid");
36 entry("sbrk");
37 entry("sleep");
38 entry("uptime");
39 entry("uniq");
40 entry("head");
41 entry("waitext");
42 entry("ps");
43
```

File: *sysproc.c*

The parameters are passed in the user program test.c

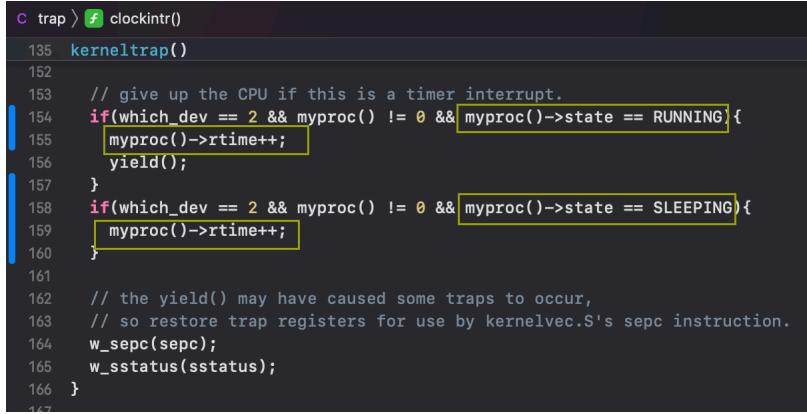


```
C sysproc > No Selection
232 uint64
233 sys_waitext(void)
234 {
235     uint64 wtime;
236     uint64 rtime;
237     uint64 p;
238
239     argaddr(0, &wtime);
240
241     argaddr(1, &rtime);
242
243     argaddr(2, &p);
244
245     return waitext(wtime, rtime, p);
246 }
247

h user > No Selection
15 int unlink(const char* );
16 int fstat(int fd, struct stat* );
17 int link(const char*, const char* );
18 int mkdir(const char* );
19 int chdir(const char* );
20 int dup(int );
21 int getpid(void );
22 char* sbrk(int );
23 int sleep(int );
24 int uptime(void );
25 int uniq(int n, char *buffer, char var_icd);
26 int head(int n, char *buffer, int lines);
27 int waitext(int *wtime, int *rtime, int *addr);
28 int ps(void );
29
```

File: *trap.c*

In *kerneltrap()* function when a timer interrupt occurs when the process state is running or sleeping, the *rtime* is incremented.



```
C trap > f clockintr()
135 kerneltrap()
152
153 // give up the CPU if this is a timer interrupt.
154 if(which_dev == 2 && myproc() != 0 && myproc()->state == RUNNING){
155     myproc()->rtime++;
156     yield();
157 }
158 if(which_dev == 2 && myproc() != 0 && myproc()->state == SLEEPING){
159     myproc()->rtime++;
160 }
161
162 // the yield() may have caused some traps to occur,
163 // so restore trap registers for use by kernelvec.S's sepc instruction.
164 w_sepc(sepc);
165 w_sstatus(sstatus);
166 }
```

File: *proc.c*

All the declared time fields are initialized in *allocproc()* function. The creation time *ctime* is incremented with ticks, as the process starts. In *exit()*, process end time *etime* is incremented with the ticks.

```

C proc > No Selection
110 allocproc(void)
148
149     // time fields initialization
150     p->ctime = ticks;           //
151     p->etime = 0;              //
152     p->rtime = 0;              //
153     p->total_time = 0;
154
155     return p;
156 }
157
```



```

C proc > No Selection
353 exit(int status)
385     p->state = ZOMBIE;
386
387     release(&wait_lock);
388
389     // updating end time
390     p->etime = ticks;
391
392     // Jump into the scheduler, never to return.
393     sched();
394     panic("zombie exit");
395 }
```

In *scheduler()* function, the *total_time* is incremented which stores the ticks whenever the process is running. In the *waitext()* function if a zombie process is found then the run time *rtime* is updated by difference of start and end time. The creation time, end time, the runtime incremented by timer interrupt, total scheduler time are printed.

```

C proc > f scheduler()
454 scheduler(void)
463
464     for(p = proc; p < &proc[NPROC]; p++) {
465         acquire(&p->lock);
466         if(p->state == RUNNABLE) {
467             //update total time
468             p->total_time += 1;
469             // Switch to chosen process. It is the
470             // to release its lock and then reacquire
471             // before jumping back to us.
472             p->state = RUNNING;
473             c->proc = p;
474             swtch(&c->context, &p->context);
475 }
```

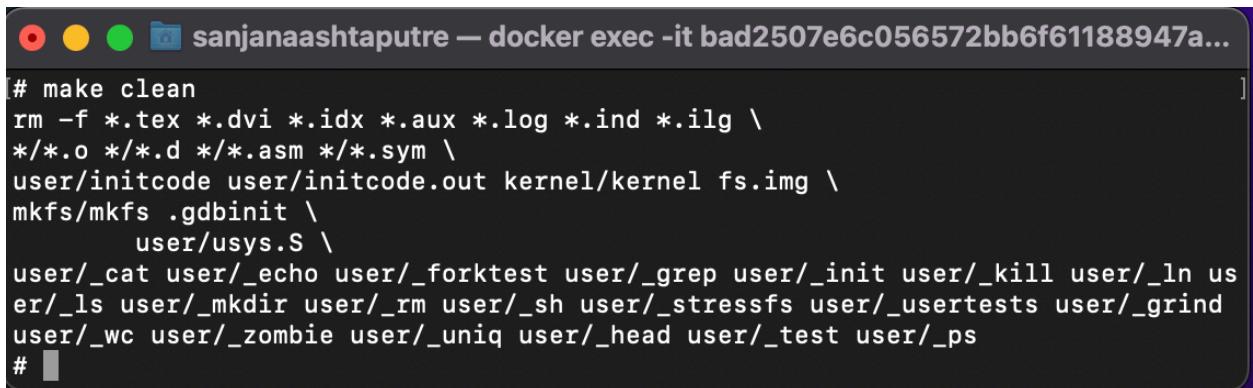


```

C proc > No Selection
697 waitext(uint64 wtime, uint64 rtime, uint64 addr)
714     if(pp->state == ZOMBIE){
715         // Found one.
716
717         //update run time
718         rtime = pp->etime - pp->ctime;
719
720         printf("\ncreation time : %d \nrun time : %d \nend time : %d \ntimer
721             interrupt while running or sleeping: %d \n", pp->ctime, rtime,
722             pp->etime, pp->rtime);
723         printf("process name : %s \ntotal scheduler time: %d \n",pp->name,
724             pp->total_time);
725 }
```

Steps to run the code and its output:

make clean



```

# make clean
rm -f *.tex *.dvi *.idx *.aux *.log *.ind *.ilg \
/*.*o */*.d /*.*.asm /*.*.sym \
user/initcode user/initcode.out kernel/kernel fs.img \
mkfs/mkfs .gdbinit \
    user/usys.S \
user/_cat user/_echo user/_forktest user/_grep user/_init user/_kill user/_ln us
er/_ls user/_mkdir user/_rm user/_sh user/_stressfs user/_usertests user/_grind
user/_wc user/_zombie user/_uniq user/_head user/_test user/_ps
# 
```

make qemu

Input files:

The image shows three separate terminal windows side-by-side. Each window has a dark background with light-colored text. The first window contains the command '\$ cat headfile.txt' followed by 20 lines of text: 'This is sentence 1.' through 'This is sentence 20.'. The second window contains the command '\$ cat headfile2.txt' followed by 8 lines of text: 'Line 1' through 'Line 8'. The third window contains the command '\$ cat uniqfile.txt' followed by 8 lines of text: 'I understand the Operating system.' repeated twice, followed by 'I love to work on OS.', 'Thanks xv6.', and 'THANKS XV6.'

```
$ cat headfile.txt
This is sentence 1.
This is sentence 2.
This is sentence 3.
This is sentence 4.
This is sentence 5.
This is sentence 6.
This is sentence 7.
This is sentence 8.
This is sentence 9.
This is sentence 10.
This is sentence 11.
This is sentence 12.
This is sentence 13.
This is sentence 14.
This is sentence 15.
This is sentence 16.
This is sentence 17.
This is sentence 18.
This is sentence 19.
This is sentence 20.$

$ cat headfile2.txt
Line 1
Line 2
Line 3
Line 4
Line 5
Line 6
Line 7
Line 8

$ cat uniqfile.txt
I understand the Operating system.
I love to work on OS.
I love to work on OS.
Thanks xv6.
THANKS XV6.
```

Output: head

test head head file.txt

The image shows a single terminal window with a dark background and light-colored text. It displays the output of the 'test' command with the argument 'head head file.txt'. The output includes the command '\$ test head headfile.txt', the process ID 'head pid 11', two sections of text labeled 'File: headfile.txt' (one in user mode, one in kernel mode), and performance statistics at the bottom: 'creation time : 2229', 'run time : 3', 'end time : 2232', 'timer interrupt while running or sleeping: 2', 'process name : head', and 'total scheduler time: 8'.

```
$ test head headfile.txt
head pid 11

File: headfile.txt
Head command is getting executed in user mode
This is sentence 1.
This is sentence 2.
This is sentence 3.
This is sentence 4.
This is sentence 5.
This is sentence 6.
This is sentence 7.
This is sentence 8.
This is sentence 9.
This is sentence 10.
This is sentence 11.
This is sentence 12.
This is sentence 13.
This is sentence 14.

File: headfile.txt
Head command is getting executed in kernel mode
This is sentence 1.
This is sentence 2.
This is sentence 3.
This is sentence 4.
This is sentence 5.
This is sentence 6.
This is sentence 7.
This is sentence 8.
This is sentence 9.
This is sentence 10.
This is sentence 11.
This is sentence 12.
This is sentence 13.
This is sentence 14.

creation time : 2229
run time : 3
end time : 2232
timer interrupt while running or sleeping: 2
process name : head
total scheduler time: 8
```

test head -10 headfile.txt

```
sanjanaashtaputre — docker exec -it bad2507e6c056572bb6f61188947a...
$ test head -10 headfile.txt
head pid 13

File: headfile.txt
Head command is getting executed in user mode
This is sentence 1.
This is sentence 2.
This is sentence 3.
This is sentence 4.
This is sentence 5.
This is sentence 6.
This is sentence 7.
This is sentence 8.
This is sentence 9.
This is sentence 10.

File: headfile.txt
Head command is getting executed in kernel mode
This is sentence 1.
This is sentence 2.
This is sentence 3.
This is sentence 4.
This is sentence 5.
This is sentence 6.
This is sentence 7.
This is sentence 8.
This is sentence 9.
This is sentence 10.

creation time : 3402
run time : 2
end time : 3404
timer interrupt while running or sleeping: 1
process name : head
total scheduler time: 3
$
```

test head -5 headfile.txt headfile2.txt

```
sanjanaashtaputre — docker exec -it bad2507e6c056572bb6f61188947a...
$ test head -5 headfile.txt headfile2.txt
head pid 21

File: headfile.txt
Head command is getting executed in user mode
This is sentence 1.
This is sentence 2.
This is sentence 3.
This is sentence 4.
This is sentence 5.

File: headfile.txt
Head command is getting executed in kernel mode
This is sentence 1.
This is sentence 2.
This is sentence 3.
This is sentence 4.
This is sentence 5.

File: headfile2.txt
Head command is getting executed in user mode
Line 1
Line 2
Line 3
Line 4
Line 5

File: headfile2.txt
Head command is getting executed in kernel mode
Line 1
Line 2
Line 3
Line 4
Line 5

creation time : 6329
run time : 2
end time : 6331
timer interrupt while running or sleeping: 3
process name : head
total scheduler time: 4
$
```

Output: uniq

test uniq uniqfile.txt

```
$ test uniq uniqfile.txt
uniq pid 25
Uniq command is getting executed in user mode.
I understand the Operating system.
I love to work on OS.
Thanks xv6.
THANKS XV6.

Uniq command is getting executed in kernel mode
I understand the Operating system.
I love to work on OS.
Thanks xv6.
THANKS XV6.

creation time : 7832
run time : 1
end time : 7833
timer interrupt while running or sleeping: 2
process name : uniq
total scheduler time: 8
$
```

test uniq -c uniqfile.txt

```
$ test uniq -c uniqfile.txt
uniq pid 27
Uniq command is getting executed in user mode.
3 I understand the Operating system.
2 I love to work on OS.
1 Thanks xv6.
1 THANKS XV6.

Uniq command is getting executed in kernel mode
3 I understand the Operating system.
2 I love to work on OS.
1 Thanks xv6.
1 THANKS XV6.

creation time : 8149
run time : 1
end time : 8150
timer interrupt while running or sleeping: 1
process name : uniq
total scheduler time: 2
$
```

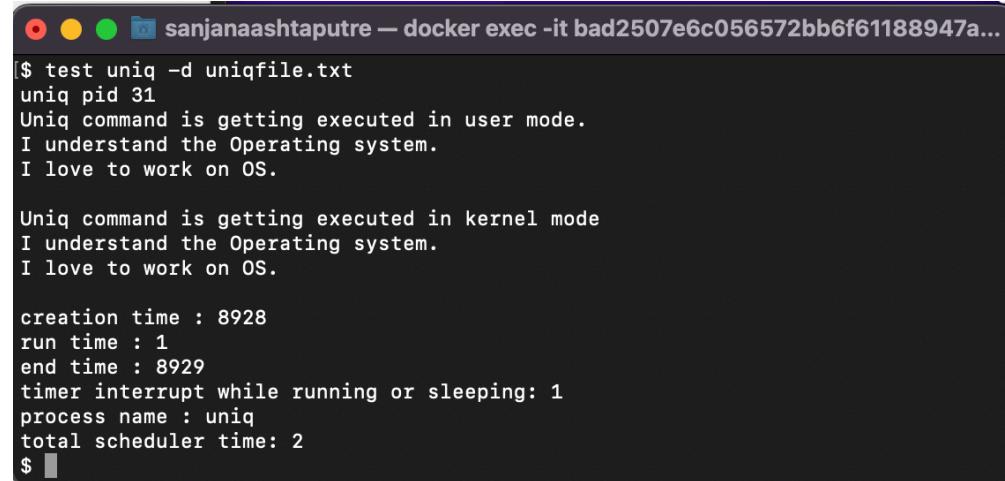
test uniq -i uniqfile.txt

```
$ test uniq -i uniqfile.txt
uniq pid 29
Uniq command is getting executed in user mode.
I understand the Operating system.
I love to work on OS.
Thanks xv6.

Uniq command is getting executed in kernel mode
I understand the Operating system.
I love to work on OS.
Thanks xv6.

creation time : 8473
run time : 0
end time : 8473
timer interrupt while running or sleeping: 0
process name : uniq
total scheduler time: 1
$
```

```
test uniq -d uniqfile.txt
```



The screenshot shows a terminal window titled "sanjanaashtaputre — docker exec -it bad2507e6c056572bb6f61188947a...". The terminal displays the following text:

```
$ test uniq -d uniqfile.txt
uniq pid 31
Uniq command is getting executed in user mode.
I understand the Operating system.
I love to work on OS.

Uniq command is getting executed in kernel mode
I understand the Operating system.
I love to work on OS.

creation time : 8928
run time : 1
end time : 8929
timer interrupt while running or sleeping: 1
process name : uniq
total scheduler time: 2
$
```