

Implementing ps on xv6

Implemented in XV6-RISCV in DOCKER (in Mac OS - M1)

Code logic:

File: *ps.c*

The system call *ps()* is called in this file which is present in the user directory of xv6.

/xv6-riscv/user/ps.c

```
1 #include "kernel/types.h"
2 #include "kernel/stat.h"
3 #include "user.h"
4
5
6 v int main(int argc, char *argv[]) {
7     ps();
8     exit(0);
9 }
10 |
```

File: *defs.h*

```
h defs > No Selection
84 // proc.c
85 int cpuid(void);
86 void exit(int);
87 int fork(void);
88 int growproc(int);
89 void proc_mapstacks(pagetable_t);
90 void proc_pagetable(struct proc *);
91 void proc_freepagetable(pagetable_t, uint64);
92 int kill(int);
93 int killed(struct proc *);
94 void setkilled(struct proc *);
95 struct cpus* mycpu(void);
96 struct cpus* getmycpu(void);
97 struct proc* myproc();
98 void procinit(void);
99 void scheduler(void) __attribute__((noreturn));
100 void sched(void);
101 void sleep(void*, struct spinlock*);
102 void userinit(void);
103 int wait(uint64);
104 void wakeup(void*);
105 void yield(void);
106 int either_copyout(int user_dst, uint64 dst, void *src, uint64 len);
107 int either_copyin(void *dst, int user_src, uint64 src, uint64 len);
108 void procdump(void);
109 int waittext(uint64, uint64, uint64);
110 int ps(void);
```

File: *proc.h*

Declared variables for process creation time (*ctime*), end time (etime), run time (*rtime*) and total time (*total_time*) in *proc* structure.

```
// Per-process state
struct proc {
    struct spinlock lock;

    // p->lock must be held when using these:
    enum proctype state;           // Process state
    void *chan;                   // If non-zero, sleeping on chan
    int killed;                   // If non-zero, have been killed
    int xstate;                   // Exit status to be returned to parent's wait
    int pid;                      // Process ID

    // wait_lock must be held when using this:
    struct proc *parent;          // Parent process

    // these are private to the process, so p->lock need not be held.
    uint64 kstack;                // Virtual address of kernel stack
    uint64 sz;                     // Size of process memory (bytes)
    pagetable_t pagetable;         // User page table
    struct trapframe *trapframe;   // data page for trampoline.S
    struct context context;        // swtch() here to run process
    struct file *ofile[NFILE];    // Open files
    struct inode * cwd;           // Current directory
    char name[16];                // Process name (debugging)

    int ctime, etime, rtime;
    uint total_time;
};
```

File: *syscall.h*

```
h syscall > No Selection
1 // System call numbers
2 #define SYS_fork 1
3 #define SYS_exit 2
4 #define SYS_wait 3
5 #define SYS_pipe 4
6 #define SYS_read 5
7 #define SYS_kill 6
8 #define SYS_exec 7
9 #define SYS_fstat 8
10 #define SYS_chdir 9
11 #define SYS_dup 10
12 #define SYS_getpid 11
13 #define SYS_sbrk 12
14 #define SYS_sleep 13
15 #define SYS_uptime 14
16 #define SYS_open 15
17 #define SYS_write 16
18 #define SYS_mknod 17
19 #define SYS_unlink 18
20 #define SYS_link 19
21 #define SYS_mkdir 20
22 #define SYS_close 21
23 #define SYS_uniq 22
24 #define SYS_head 23
25 #define SYS_waittext 24
26 #define SYS_ps 25
27
```

File: *syscall.c*

```
c syscall > No Selection
103 extern uint64 sys_exit(void);
104 extern uint64 sys_uniq(void);
105 extern uint64 sys_head(void);
106 extern uint64 sys_waittext(void);
107 extern uint64 sys_ps(void);
108
109 // An array mapping syscall numbers from syscall.h
110 // to the function that handles the system call.
111 static uint64 (*syscalls[])() = {
112 [SYS_fork] sys_fork,
113 [SYS_exit] sys_exit,
114 [SYS_wait] sys_wait,
115 [SYS_pipe] sys_pipe,
116 [SYS_read] sys_read,
117 [SYS_kill] sys_kill,
118 [SYS_exec] sys_exec,
119 [SYS_fstat] sys_fstat,
120 [SYS_chdir] sys_chdir,
121 [SYS_dup] sys_dup,
122 [SYS_getpid] sys_getpid,
123 [SYS_sbrk] sys_sbrk,
124 [SYS_sleep] sys_sleep,
125 [SYS_uptime] sys_uptime,
126 [SYS_open] sys_open,
127 [SYS_write] sys_write,
128 [SYS_mknod] sys_mknod,
129 [SYS_unlink] sys_unlink,
130 [SYS_link] sys_link,
131 [SYS_mkdir] sys_mkdir,
132 [SYS_close] sys_close,
133 [SYS_uniq] sys_uniq,
134 [SYS_head] sys_head,
135 [SYS_waittext] sys_waittext,
136 [SYS_ps] sys_ps,
137 };
138
```

File: *usys.pl*

```
☒ usys > No Selection
25 entry("Kill");
26 entry("exec");
27 entry("open");
28 entry("mknod");
29 entry("unlink");
30 entry("fstat");
31 entry("link");
32 entry("mkdir");
33 entry("chdir");
34 entry("dup");
35 entry("getpid");
36 entry("sbrk");
37 entry("sleep");
38 entry("uptime");
39 entry("uniq");
40 entry("head");
41 entry("waittext");
42 entry("ps");
43
```

File: *sysproc.c*

```
C sysproc > No Selection
234 sys_waittext(void)
235     uint64 rtime;
236     uint64 p;
237
238     argaddr(0, &wtime);
239
240     argaddr(1, &rtime);
241
242     argaddr(2, &p);
243
244     return waittext(wtime, rtime, p);
245 }
246
247 int
248 sys_ps(void)
249 {
250     return ps();
251 }
```

File: *user.h*

```
h user > No Selection
17 int link(const char*, const char*);
18 int mkdir(const char*);
19 int chdir(const char*);
20 int dup(int);
21 int getpid(void);
22 char* sbrk(int);
23 int sleep(int);
24 int uptime(void);
25 int uniq(int n, char *buffer, char var_icd);
26 int head(int n, char *buffer, int lines);
27 int waittext(int *wtime, int *rtime, int *addr);
28 int ps(void);
29
```

File: *proc.c*

```
C proc > No Selection
110 allocproc(void)
111
112     // time fields initialization
113     p->ctime = ticks;           //
114     p->etime = 0;              //
115     p->rtime = 0;              //
116     p->total_time = 0;
117
118     return p;
119 }
```

```
C proc > f scheduler()
454 scheduler(void)
455
456     for(p = proc; p < &proc[NPROC]; p++) {
457         acquire(&p->lock);
458         if(p->state == RUNNABLE) {
459             //update total time
460             p->total_time += 1;
461
462             // Switch to chosen process. It is the
463             // to release its lock and then reacqur
464             // before jumping back to us.
465             p->state = RUNNING;
466             c->proc = p;
467             swtch(&c->context, &p->context);
468     }
```

All the declared time fields are initialized in *allocproc()* function. The creation time is incremented with ticks, as the process starts. In *scheduler()* function, the *total_time* is incremented which stores the ticks whenever the process is running. In *ps()* function, the process name, pid, start time, process status, total time is displayed for unused (previous process that is now stored as unused), used, sleeping, running, runnable and zombie processes.

```
c proc > No Selection
751 int
752 ps()
753 {
754     struct proc *p;
755
756     //Loop over process table looking for process with pid.
757     acquire(&wait_lock);
758     printf("Name \t PID \t Status \t Start_time \t Total Scheduler time\n");
759     for(p = proc; p < &proc[NPROC]; p++){
760         if(p->state == UNUSED && p->ctime != 0){
761             printf("%s \t %d \t UNUSED ", p->name,p->pid);
762             printf(" \t %d ",p->ctime);
763             printf(" \t \t %d \n",p->total_time);
764         }
765
766         else if(p->state == USED){
767             printf("%s \t %d \t USED ", p->name,p->pid);
768             printf(" \t %d ",p->ctime);
769             printf(" \t \t %d \n",p->total_time);
770         }
771
772         else if(p->state == SLEEPING){
773             printf("%s \t %d \t SLEEPING ", p->name,p->pid);
774             printf(" \t %d ",p->ctime);
775             printf(" \t \t %d \n",p->total_time);
776         }
777         else if(p->state == RUNNING){
778             printf("%s \t %d \t RUNNING ", p->name,p->pid);
779             printf(" \t %d ",p->ctime);
780             printf(" \t \t %d \n",p->total_time);
781         }
782         else if(p->state == RUNNABLE){
783             printf("%s \t %d \t RUNNABLE ", p->name,p->pid);
784             printf(" \t %d ",p->ctime);
785             printf(" \t \t %d \n",p->total_time);
786         }
787         else if(p->state == ZOMBIE){
788             printf("%s \t %d \t ZOMBIE ", p->name,p->pid);
789             printf(" \t %d ",p->ctime);
790             printf(" \t \t %d \n",p->total_time);
791         }
792     }
793     release(&wait_lock);
794     return 25;
795 }
796
```

Steps to run the code and its output

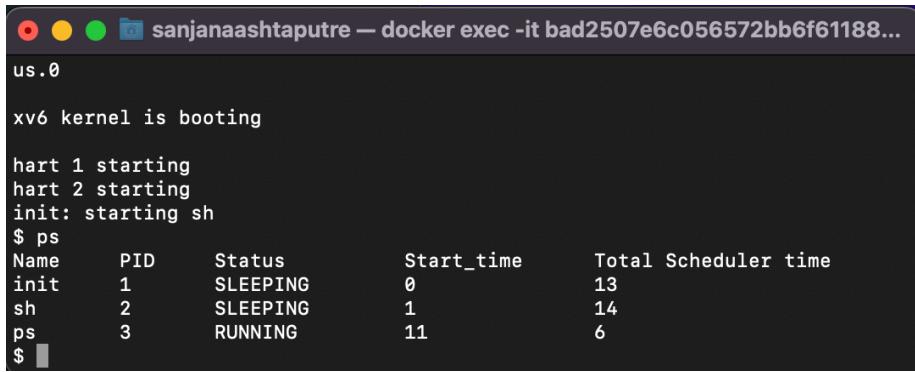
make clean

```
sanjanaashtaputre — docker exec -it bad2507e6c056572bb6f61188947a...
[ # make clean
rm -f *.tex *.dvi *.idx *.aux *.log *.ind *.ilg \
/*.* /*.*.d /*.*.asm /*.*.sym \
user/initcode user/initcode.out kernel/kernel fs.img \
mkfs/mkfs .gdbinit \
        user/usys.S \
user/_cat user/_echo user/_forktest user/_grep user/_init user/_kill user/_ln us
er/_ls user/_mkdir user/_rm user/_sh user/_stressfs user/_usertests user/_grind
user/_wc user/_zombie user/_uniq user/_head user/_test user/_ps
# ]
```

make qemu

Output:

ps as the first command with no previous user command.

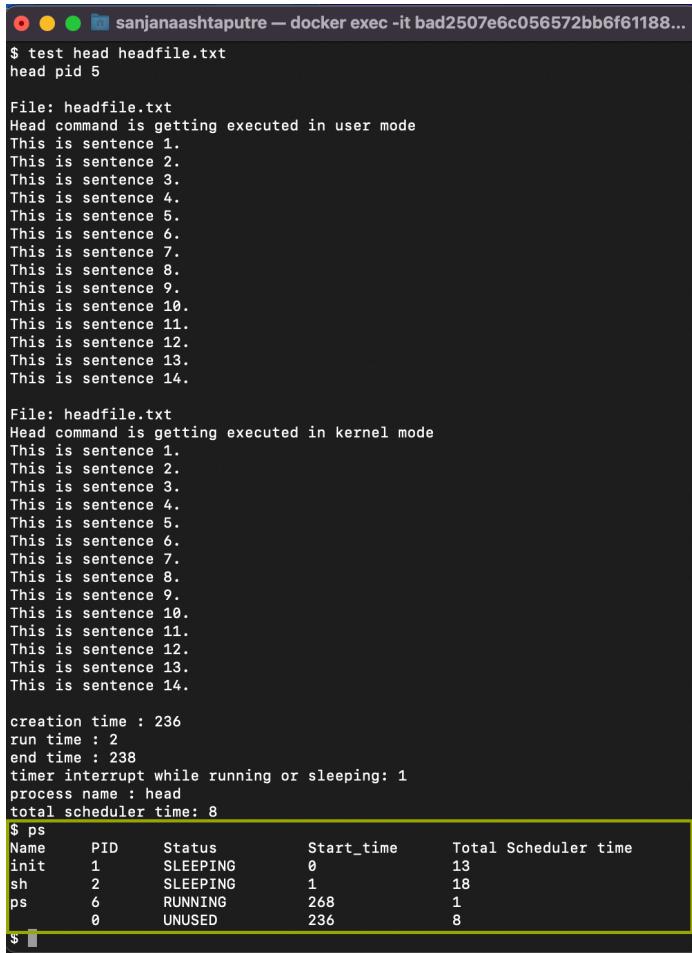


```
sanjanaashtaputre — docker exec -it bad2507e6c056572bb6f61188...
us.0

xv6 kernel is booting

hart 1 starting
hart 2 starting
init: starting sh
$ ps
Name      PID      Status        Start_time      Total Scheduler time
init      1      SLEEPING      0              13
sh        2      SLEEPING      1              14
ps        3      RUNNING       11             6
$
```

ps after a command is executed. On successful execution, the process is stored with unused status.



```
sanjanaashtaputre — docker exec -it bad2507e6c056572bb6f61188...
$ test head headfile.txt
head pid 5

File: headfile.txt
Head command is getting executed in user mode
This is sentence 1.
This is sentence 2.
This is sentence 3.
This is sentence 4.
This is sentence 5.
This is sentence 6.
This is sentence 7.
This is sentence 8.
This is sentence 9.
This is sentence 10.
This is sentence 11.
This is sentence 12.
This is sentence 13.
This is sentence 14.

File: headfile.txt
Head command is getting executed in kernel mode
This is sentence 1.
This is sentence 2.
This is sentence 3.
This is sentence 4.
This is sentence 5.
This is sentence 6.
This is sentence 7.
This is sentence 8.
This is sentence 9.
This is sentence 10.
This is sentence 11.
This is sentence 12.
This is sentence 13.
This is sentence 14.

creation time : 236
run time : 2
end time : 238
timer interrupt while running or sleeping: 1
process name : head
total scheduler time: 8
$ ps
Name      PID      Status        Start_time      Total Scheduler time
init      1      SLEEPING      0              13
sh        2      SLEEPING      1              18
ps        6      RUNNING       268             1
                  0      UNUSED        236             8
$
```

When `ps()` is added in the test.c file.

```
/xv6-riscv/user/test.c

11 v     for(int i=1; i<argc; i++){
12         argv2[i-1] = argv[i];
13     }
14     ps();
15
16     pid = fork();
17 v     if (pid == 0) {
18         // Child process
19         exec(argv[1], argv2);
20         printf("Failed to execute %s\n", argv[1]);
21         exit(1);
22 v     } else if (pid > 0) {
23         printf("%s pid %d \n", argv[1], pid);
24         // Parent process
25         waitext(&wt, &rt, &st);
26 v     } else {
27         printf("Failed to fork %s\n", argv[1]);
28     }
29     ps();
30
31     exit(0);
32 }
33
```

```
sanjanaashtaputre — docker exec -it bad2507e6c056572bb6f61188...
$ test head -5 headfile.txt
Name      PID      Status       Start_time      Total Scheduler time
init      1        SLEEPING    0              25
sh        2        SLEEPING    2              13
test      3        RUNNING    77             6
head pid 4

File: headfile.txt
Head command is getting executed in user mode
This is sentence 1.
This is sentence 2.
This is sentence 3.
This is sentence 4.
This is sentence 5.

File: headfile.txt
Head command is getting executed in kernel mode
This is sentence 1.
This is sentence 2.
This is sentence 3.
This is sentence 4.
This is sentence 5.

creation time : 78
run time : 1
end time : 79
timer interrupt while running or sleeping: 1
process name : head
total scheduler time: 8
Name      PID      Status       Start_time      Total Scheduler time
init      1        SLEEPING    0              25
sh        2        SLEEPING    2              13
test      3        RUNNING    77             7
          0        UNUSED      78             8
```